- snowflake-ml-python version: 1.2.2
- Feature Store PrPr Version: 0.5.1
- Updated date: 2/12/2024

Basic Feature Demo

This notebook demonstrates feature store with simple features. It includes an end-2-end ML experiment cycle: feature creation, training and inference. It also demonstrate the interoperation between Feature Store and Model Registry.

```
In []: from snowflake.snowpark import Session
    from snowflake.snowpark import functions as F
    from snowflake.ml.feature_store import (
        FeatureStore,
        FeatureView,
        Entity,
        CreationMode
)
    from snowflake.ml.utils.connection_params import SnowflakeLoginOptions
```

Setup Snowflake connection and database

For detailed session connection config, please follow this tutorial.

```
In [ ]: session = Session.builder.configs(SnowflakeLoginOptions()).create()
```

Below cell creates temporary database, schema and warehouse for this notebook. All temporary resources will be deleted at the end of this notebook. You can rename with your own name if needed.

```
{FS_DEMO_DB}.{TEST_DATASET_SCHEMA}""").collect()
session.sql(f"CREATE WAREHOUSE IF NOT EXISTS {FS_DEMO_WH}").collect()
```

Create FeatureStore Client

Let's first create a feature store client.

We can pass in an existing database name, or a new database will be created upon the feature store initialization. Replace DEMO_DB and DEMO_SCHEMA with your database and schema.

Prepare test data

We will use wine quality dataset for this demo. Download the public dataset from kaggle if you dont have it already: https://www.kaggle.com/datasets/uciml/red-wine-quality-cortez-et-al-2009. Replace TEST_CSV_FILE_PATH with your local file path.

```
In [ ]: TEST_CSV_FILE_PATH = 'winequality-red.csv'
        session.file.put(
            f"file://{TEST_CSV_FILE_PATH}", session.get_session_stage())
        from snowflake.snowpark.types import (
            StructType,
            StructField,
            IntegerType,
            FloatType
        input_schema = StructType(
                StructField("fixed_acidity", FloatType()),
                StructField("volatile_acidity", FloatType()),
                StructField("citric_acid", FloatType()),
                StructField("residual_sugar", FloatType()),
                StructField("chlorides", FloatType()),
                StructField("free_sulfur_dioxide", IntegerType()),
                StructField("total_sulfur_dioxide", IntegerType()),
                StructField("density", FloatType()),
                StructField("pH", FloatType()),
                StructField("sulphates", FloatType()),
                StructField("alcohol", FloatType()),
                StructField("quality", IntegerType())
```

Create and register a new Entity

We will create an Entity called wine and register it with the feature store.

You can retrieve the active Entities in the feature store with list_entities() API.

```
In []: entity = Entity(name="WINE", join_keys=["WINE_ID"])
    fs.register_entity(entity)
    fs.list_entities().show()
```

Load source data and do some simple feature engineering

Then we will load from the source table and conduct some simple feature engineerings.

Here we are just doing two simple data manipulation (but more complex ones are carried out the same way):

- 1. Assign a WINE_ID column to the source
- 2. Derive a new column by multipying two existing feature columns

```
In [ ]: source_df = session.table(full_table_name)
        source_df.show()
In [ ]: def addIdColumn(df, id_column_name):
            # Add id column to dataframe
            columns = df.columns
            new df = df.withColumn(id column name, F.monotonically increasing id())
            return new_df[[id_column_name] + columns]
        def generate new feature(df):
            # Derive a new feature column
            return df.withColumn(
                "MY NEW FEATURE", df["FIXED ACIDITY"] * df["CITRIC ACID"])
        df = addIdColumn(source_df, "WINE_ID")
        feature df = generate new feature(df)
        feature_df = feature_df.select(
                'WINE_ID',
                'FIXED_ACIDITY',
                'VOLATILE_ACIDITY',
                'CITRIC_ACID',
                'RESIDUAL SUGAR',
```

```
'CHLORIDES',
'FREE_SULFUR_DIOXIDE',
'TOTAL_SULFUR_DIOXIDE',
'DENSITY',
'PH',
'MY_NEW_FEATURE',
]
)
feature_df.show()
```

Create a new FeatureView and materialize the feature pipeline

Once the FeatureView construction is done, we can materialize the FeatureView to the Snowflake backend and incremental maintenance will start.

```
In [ ]: # NOTE:
        # Due to a known issue on backend pipeline creation,
        # if the source data is created right before the
        # feature pipeline, there might be a chance for
        # dataloss, so sleep for 60s for now.
        # This issue will be fixed soon in upcoming patch.
        import time
        time.sleep(60)
In [ ]: fv = FeatureView(
            name="WINE FEATURES",
            entities=[entity],
            feature_df=feature_df,
            refresh_freq="1 minute",
            desc="wine features"
        fv = fs.register_feature_view(
            feature_view=fv,
            version="V1",
            block=True
In [ ]: # Examine the FeatureView content
        fs.read_feature_view(fv).show()
```

Explore additional features

Now I have my FeatureView created with a collection of features, but what if I want to explore additional features on top?

Since a materialized FeatureView is immutable (due to singe DDL for the backend dynamic table), we will need to create a new FeatureView for the additional features and then merge them.

```
In [ ]: extra_feature_df = df.select(
                 'WINE ID',
                'SULPHATES',
                'ALCOHOL',
        new fv = FeatureView(
            name="EXTRA_WINE_FEATURES",
            entities=[entity],
            feature_df=extra_feature_df,
            refresh_freq="1 minute",
            desc="extra wine features"
        new_fv = fs.register_feature_view(
            feature_view=new_fv,
            version="V1",
            block=True
In [ ]: | # We can easily retrieve all FeatureViews for a given Entity.
        fs.list feature views(entity name="WINE"). \
            select(["NAME", "ENTITIES", "FEATURE_DESC"]).show()
```

Create new feature view with combined feature results [Optional]

Now we have two FeatureViews ready, we can choose to create a new one by merging the two (it's just like a join and we provide a handy function for that). The new FeatureView won't incur the cost of feature pipelines but only the table join cost.

Obviously we can also just work with two separate FeatureViews (most of our APIs support multiple FeatureViews), the capability of merging is just to make the features better organized and easier to share.

Generate Training Data

After our feature pipelines are fully setup, we can start using them to generate training data and later do model prediction.

```
In []: spine_df = session.table(f"{FS_DEMO_DB}.{TEST_DATASET_SCHEMA}.WINE_DATA")
    spine_df = addIdColumn(source_df, "WINE_ID")
    spine_df = spine_df.select("WINE_ID", "QUALITY")
    spine_df.show()
```

Generate training data is easy since materialized FeatureViews already carry most of the metadata like join keys, timestamp for point-in-time lookup, etc. We just need to provide the spine data (it's called spine because we are essentially enriching the data by joining features with it). We can also generate dataset with a subset of features in the feature view by slice.

```
In [ ]: training_dataset_full_path = \
            f"{FS DEMO DB}.{FS DEMO SCHEMA}.WINE TRAINING DATA TABLE"
        session.sql(f"DROP TABLE IF EXISTS {training_dataset_full_path}") \
            .collect()
        training data = fs.generate dataset(
            spine_df=spine_df,
            features=[
                full fv.slice([
                    "FIXED_ACIDITY",
                    "VOLATILE ACIDITY",
                    "CITRIC ACID"
                ])
            ],
            materialized_table="WINE_TRAINING_DATA_TABLE",
            spine_timestamp_col=None,
            spine_label_cols=["QUALITY"],
            save mode="merge",
            exclude columns=['WINE ID']
        training data.df.show()
```

Train model with Snowpark ML

Now let's training a simple random forest model, and evaluate the prediction accuracy.

```
In [ ]: from snowflake.ml.modeling.ensemble import RandomForestRegressor
        from snowflake.ml.modeling import metrics as snowml metrics
        from snowflake.snowpark.functions import abs as sp_abs, mean, col
        def train model using snowpark ml(training data):
            train, test = training data.df.random split([0.8, 0.2], seed=42)
            feature_columns = [col for col in training_data.df.columns if col != "QU
            label column = "QUALITY"
            rf = RandomForestRegressor(
                input cols=feature columns, label cols=[label column],
                max depth=3, n estimators=20, random state=42
            )
            rf.fit(train)
            predictions = rf.predict(test)
            mse = snowml metrics.mean squared error(
                df=predictions,
                y_true_col_names=label_column,
```

```
y_pred_col_names="OUTPUT_" + label_column)

accuracy = 100 - snowml_metrics.mean_absolute_percentage_error(
    df=predictions,
    y_true_col_names=label_column,
    y_pred_col_names="OUTPUT_" + label_column
)

print(f"MSE: {mse}, Accuracy: {accuracy}")
    return rf

rf = train_model_using_snowpark_ml(training_data)
print(rf)
```

[Predict Optional 1] With local model

Now let's predict with the model and the feature values retrieved from feature store.

[Predict Option 2] Using Model Registry

Step 1: Log the model along with its training dataset metadata into Model Registry

```
In []: from snowflake.ml.registry import model_registry

registry = model_registry.ModelRegistry(
    session=session,
    database_name=MR_DEMO_DB,
    create_if_not_exists=True
)
```

Register the dataset into model registry with <code>log_artifact</code>. Artifact is a generalized concept of ML pipeline outputs that are needed for subsequent execution. Refer to https://docs.snowflake.com/LIMITEDACCESS/snowflake-ml-model-registry for more details about the API.

```
In [ ]: DATASET_NAME = "MY_DATASET"
DATASET_VERSION = "V1"

my_dataset = registry.log_artifact(
```

```
artifact=training_data,
name=DATASET_NAME,
version=DATASET_VERSION,
)
```

Now you can log the model together with the registered artifact (which is a dataset here).

```
In []: model_name = "MY_MODEL"

model_ref = registry.log_model(
    model_name=model_name,
    model_version="V2",
    model=rf,
    tags={"author": "my_rf_with_training_data"},
    artifacts=[my_dataset],
    options={"embed_local_ml_library": True},
)
```

Step 2: Restore model and predict with features

We retrieve the training dataset from registry then construct dataframe of latest feature values. Then we restore the model from registry. At last, we can predict with latest feature values.

Cleanup notebook

Cleanup resources created in this notebook.

```
In [ ]: session.sql(f"DROP DATABASE IF EXISTS {FS_DEMO_DB}").collect()
```

```
session.sql(f"DROP DATABASE IF EXISTS {MR_DEMO_DB}").collect()
session.sql(f"DROP WAREHOUSE IF EXISTS {FS_DEMO_WH}").collect()
```